

50.017 Graphics and Visualization

Assignment 4 – Ray Tracing

Handout date: 2024.03.21

Submission deadline: 2024.04.01, 11:59 pm

Late submissions are not accepted




Figure 1. From left to right and then top to bottom: expected rendering outputs of scene 1-6.

In this assignment, you will implement a basic ray tracer. As seen in lecture, a ray tracer sends a ray for each pixel and intersects it with all the objects in the scene. Your ray tracer will support perspective cameras as well as several primitives (spheres, planes, and triangles). You will also have to support Phong lighting.

1. User Interface

In this assignment, there is no longer an interactive window for manipulating and rendering 3D models. Instead, the program takes a scene file that describes the camera, lights and objects information as input. The output is a rendered image saved at a user-specified path.

Your program should take a number of command line arguments to specify the input file, output image size, and output file. To run your program on all the six test cases, you can type commands as follows:

```
YourProgramName -input scene01_plane.txt -size 200 200 -output scene_1.bmp
YourProgramName -input scene02_cube.txt -size 200 200 -output scene_2.bmp
YourProgramName -input scene03_sphere.txt -size 200 200 -output scene_3.bmp
YourProgramName -input scene04_axes.txt -size 200 200 -output scene_4.bmp
YourProgramName -input scene05_bunny_1k.txt -size 200 200 -output scene_5.bmp
YourProgramName -input scene06_shine.txt -size 200 200 -output scene_6.bmp
```

2. Basic Ray Tracer

You will use object-oriented design to make your ray tracer flexible and extendable. A generic `Object3D` class will serve as the parent class for all 3D primitives. You will derive subclasses (such as `Sphere`, `Plane`, `Triangle`, `Group`, `Transform`, `Mesh`) to implement specialized primitives. Also, this assignment requires the implementation of a `material` class for Phong lighting computation.

2.1 Ray Tracing Pipeline

Main. The `main` function in `main.cpp` reads the scene (using the `SceneParse` class), loops over the pixels in the image plane, generates a ray using your `camera` class, intersects it with the high-level `Group` that stores the objects of the scene, and writes the color of the closest intersected object.

Note that `SceneParse` is completely implemented for you. Use it to load the camera, background color and objects of the scene from scene files.

2.2 Ray Generation

Camera. In `PerspectiveCamera` class that inherits `Camera`, the field of view is specified with an angle (as shown in Figure 2). Here *up* and *direction* are not necessarily perpendicular. The \mathbf{u} , \mathbf{v} , \mathbf{w} vectors are computed using cross products. $\mathbf{w} = \text{direction}$, $\mathbf{u} = \mathbf{w} \times \text{up}$, $\mathbf{v} = \mathbf{u} \times \mathbf{w}$. The camera does not know about screen resolution. Image resolution should be handled in your main loop.

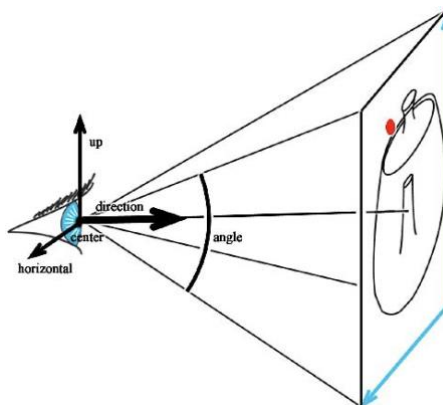


Figure 2. Perspective virtual camera.

The `Camera` class has two pure virtual methods:

```
virtual Ray generateRay( const Vector2f& point ) = 0;
virtual float getTMin() const = 0;
```

The first is used to generate rays for each screen-space coordinate, described as a `Vector2f`. The `getTMin()` method will be useful when tracing rays through the scene. For a perspective camera, the value of `tmin` will be zero to correctly clip objects behind the viewpoint (already implemented for you).

2.3 Scene Objects

Object3D. Look at the virtual `Object3D` class. It only provides the specification for 3D primitives, and in particular the ability to be intersected with a ray via the virtual method:

```
virtual bool intersect(const Ray& r, Hit& h, float tmin) = 0;
```

Since this method is pure virtual for the `Object3D` class, the prototype in the header file includes `'= 0'`. This `'= 0'` tells the compiler that `Object3D` won't implement the method, but that subclasses derived from `Object3D` must implement this routine. An `Object3D` stores a pointer to its `Material` type. The `Object3D` class has:

- a default constructor and destructor
- a pointer to a `Material` instance
- a pure virtual intersection method

Group. Fill in `Group`, also a subclass of `Object3D`, that stores a list of pointers to `Object3D` instances. For example, it will be used to store all objects in the entire scene. The `bool intersect(const Ray& r , Hit& h , float tmin)` function of `Group` loops through all these instances, calling their intersection methods. The `Group` constructor should take as input the number of objects under the group. The group also includes a method to add objects.

3. Your Tasks

3.1 Ray Intersection

Sphere. Fill in `Sphere`, a subclass of `Object3D`, that additionally stores a center point and a radius. The `Sphere` constructor will be given a center, a radius, and a pointer to a `Material` instance. The `Sphere` class implements the virtual `intersect` method mentioned above (but without the `'= 0'`):

```
virtual bool intersect(const Ray& r, Hit& h, float tmin);
```

With the `intersect` routine, we are looking for the closest intersection along a `Ray`, parameterized by `t`. `tmin` is used to restrict the range of intersection. If an intersection is found such that `t > tmin` and `t` is less than the value of the intersection currently stored in the `Hit` data structure, `Hit` is updated as necessary. Note that if the new intersection is closer than the previous one, `t`, `Material`, and `Normal` must be modified. It is important that your intersection routine verifies that `t >= tmin`. Note that `tmin` is not modified by the intersection routine.

Plane. Implement `Plane`, an infinite plane primitive derived from `Object3D`. The constructor is assumed to be:

```
Plane( const Vector3f& normal, float d, Material* m );
```

d is the offset from the origin, meaning that the plane equation is $\mathbf{P} \cdot \mathbf{n} = d$.

Implement the virtual `intersect` method mentioned above (but without the `'= 0'`), and remember that you also need to update the normal stored by `Hit`, in addition to the intersection distance and color.

```
virtual bool intersect(const Ray& r, Hit& h, float tmin);
```

Triangle. Fill in triangle primitive which also derives from `Object3D`. The constructor takes 3 vertices:

```
Triangle( const Vector3f& a, const Vector3f& b, const Vector3f& c,  
Material* m );
```

Use the method of your choice to implement the ray-triangle intersection, preferably using barycentric coordinates. Suppose we have barycentric coordinates $\lambda_0, \lambda_1, \lambda_2$ and vertex normals $\mathbf{n}_0, \mathbf{n}_1, \mathbf{n}_2$, the interpolated normal can be computed as $\lambda_0\mathbf{n}_0 + \lambda_1\mathbf{n}_1 + \lambda_2\mathbf{n}_2$.

Transform. Fill in subclass `Transform` from `Object3D`. Similar to a `Group`, a `Transform` will store a pointer to an `Object3D` (but only one, not an array). The constructor of a `Transform` takes a 4×4 matrix as input and a pointer to the `Object3D` modified by the transformation:

```
Transform( const Matrix4f& m, Object3D* o );
```

The `intersect` routine will first transform the ray, then delegate to the `intersect` routine of the contained object. Make sure to correctly transform the resulting normal according to this rule: if we transform a vertex by a homogeneous transformation matrix \mathbf{M} , its normal should be transformed by the matrix $(\mathbf{M}^{-1})^T$. A discussion of why this is the case appears in this wiki page: [https://en.wikipedia.org/wiki/Normal_\(geometry\)](https://en.wikipedia.org/wiki/Normal_(geometry)) (see Transforming Normals section).

You may choose to normalize the direction of the transformed ray or leave it un-normalized. If you decide not to normalize the direction, you might need to update some of your intersection code.

3.2 Lighting Computations

Material. Implement the diffuse and specular components of Phong lighting in `Material` class. You may refer to Lecture 11 for the equations of these two components.

```
Vector3f Shade( const Ray& ray, const Hit& hit, const Vector3f&  
dirToLight, const Vector3f& lightColor )
```

Light. We provide the pure virtual `Light` class and two subclasses: directional light and point light. Scene lighting can be accessed with

```
SceneParser::getLight() and  
SceneParser::getAmbientLight() methods.
```

Use the `Light` method:

```
void getIllumination( const Vector3f& p, Vector3f& dir, Vector3f&  
col );
```

to find the illumination at a particular location in space. p is the intersection point that you want to shade, and the function returns the normalized direction toward the light source in `dir` and the light color and intensity in `col`.

3. Hints

- Incremental debugging. Implement and test one primitive at a time. Test one lighting component at a time: ambient, diffuse, and specular.
- Use a small image size for faster debugging. 64×64 pixels is usually enough to realize that something might be wrong.
- As usual, don't hesitate to print as much information as needed for debugging, such as the direction vector of the rays, the hit values, etc.
- Use `assert()` to check function preconditions, array indices, etc. See `cassert`.
- The "very large" negative and positive values for t used in the Hit class and the intersect routine can simply be initialized with large values relative to the camera position and scene dimensions. However, to be more correct, you can use the positive and negative values for infinity from the IEEE floating point standard.
- Parse the arguments of the program in a separate function. It will make your code easier to read.
- Implement the normal visualization and Phong lighting before the transformations.
- Use the various rendering modes (normal, diffuse, distance) to debug your code. This helps you locate which part of your code is buggy.

4. Grading

Each part of this assignment is weighted as follows:

- Ray Intersection: 70%
 - Sphere: 15%
 - Plane: 15%
 - Triangle: 20%
 - Transform: 20%
- Lighting Computations: 30%

5. Submission

A .zip compressed file renamed to AssignmentN_Name_I.zip, where N is the number of the current assignment, Name is your first name, and I is the number of your student ID. It should contain only:

- The **source code** project folder (the entire thing).
- A **readme.txt** file containing a description of how you solved each part of the assignment (use the same titles) and whatever problems you encountered. If you know there are bugs in your code, please provide a list of them, and describe what do you think caused it if possible. This is very important as we can give you partial credit if you help us understand your code better.
- Six **screenshots** clearly showing that you can render images of the six test scenes with ray tracing.

Upload your zipped assignment to e-dimension. Late submissions receive 0 points!